

TCP Congestion Control -- Overview

Kameswari Chebrolu

Seminal Paper: Congestion Avoidance and Control
by Van Jacobson and Michael J. Karels

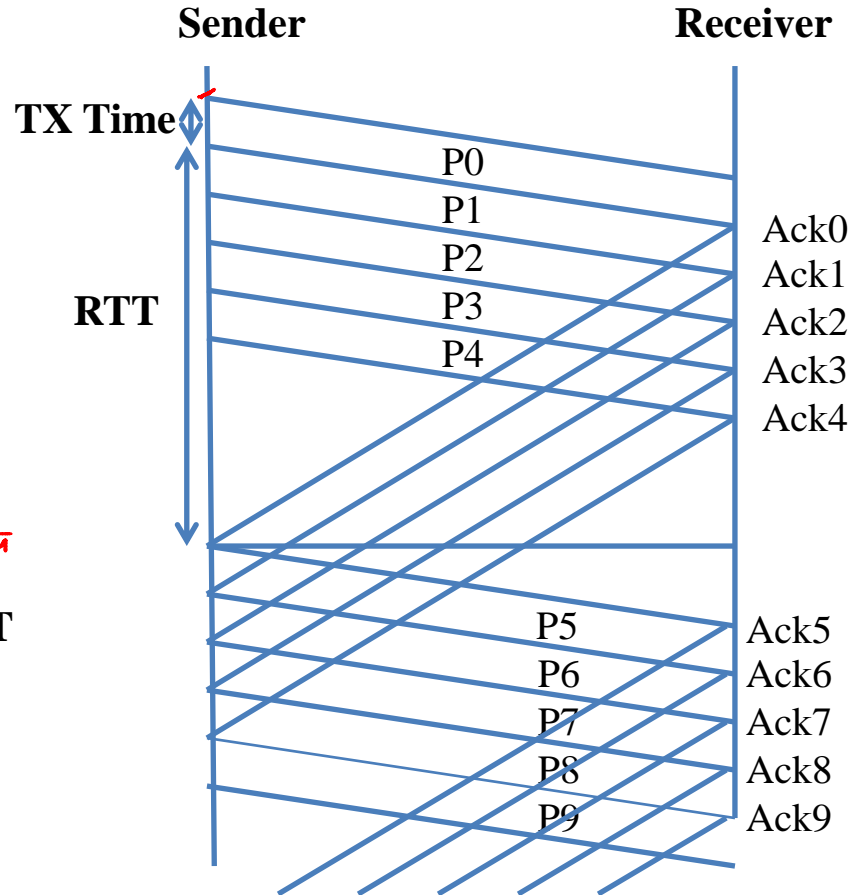
All the figures used as part of the slides are either self created or from the public domain with either 'creative commons' or 'public domain dedication' licensing. The public sites from which some of the figures have been picked include:
<http://commons.wikimedia.org> (Wikipedia, Wikimedia and workbooks); <http://www.sxc.hu> and <http://www.pixabay.com>

Recap: TCP Services

- Multiplexing/Demultiplexing
- Reliable point-to-point data transfer
- Full-duplex
- Congestion control
- Flow control

Sliding
window
protocol

Recap: Sliding Window




Throughput $\sim (W * MSS) / RTT$

Handwritten annotations:


- A red circle around **W** with an arrow pointing to it and the text "wind. size" (window size).
- A red arrow pointing up to **MSS** with the text "pkts" (packets).

Congestion Control: Problem Statement



- Many flows pass through a router; number varies with time
- Flows can be TCP or UDP ^{fairness}
- The link capacities of the routers are different
- End Result: Throughput achieved by a given flow function of many factors ^{blocks}

Congestion Control: Challenge




- Need to estimate W (of sliding window) such that each flow gets its fair share
 - Estimate small \rightarrow underutilization; Estimate large \rightarrow Congestion
- W will vary over time
- Congestion Control: Preventing sources from sending too much data too fast and thereby ‘congest’ the network

Sliding Window Protocol

- Roughly, idea translates to the following:
- View network as a pipe
- Determine the capacity of the pipe (Bandwidth-delay product)
- Fill the pipe with data
- As you remove one packet from the pipe, add another
 - ACKs help clock out data (Self Clocking)

Self-Clocking



W → window size in bits

① → $W = B \times RTT$

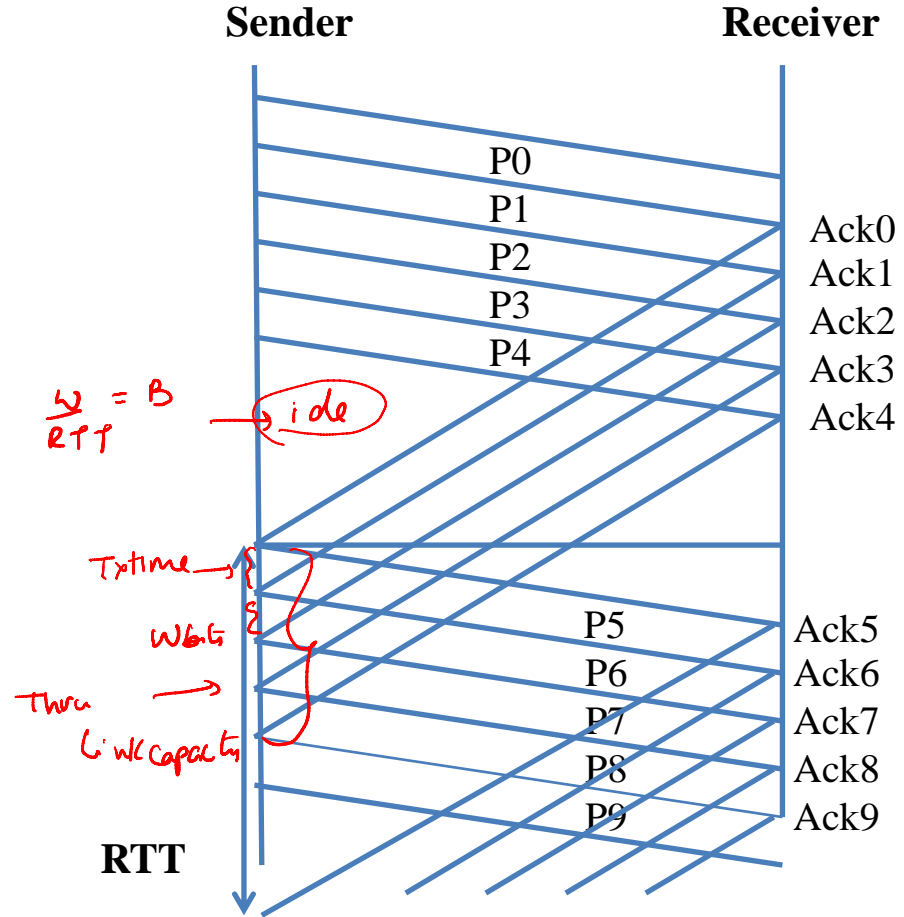
② Max Thr A & B - ? = B

③ Rate @ which Ack's are coming?
 → Rate @ which Pkts
 ↳ Link Capacity

$W < B \times RTT$

→ Link capacity

Self-clocking



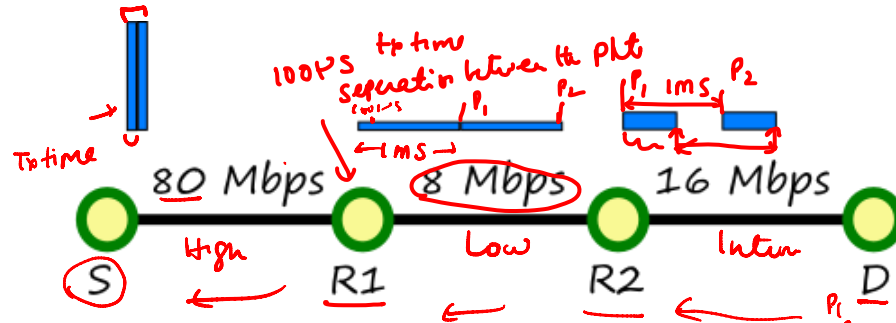

Self-Clocking

RTT \times 100 Kbps

① $w = ?$

② Max Thr. = ?

③ rate @ acks @ sender
 \hookrightarrow rate @ pkts on this bottleneck link



1KB TCP
Data pkts

TCP ACKs

RTT

Receiver

P0

P1

P2

P3

P4

Ack0

Ack1

Ack2

Ack3

Ack4

P5

P6

P7

P8

P9


Ack5

Ack6

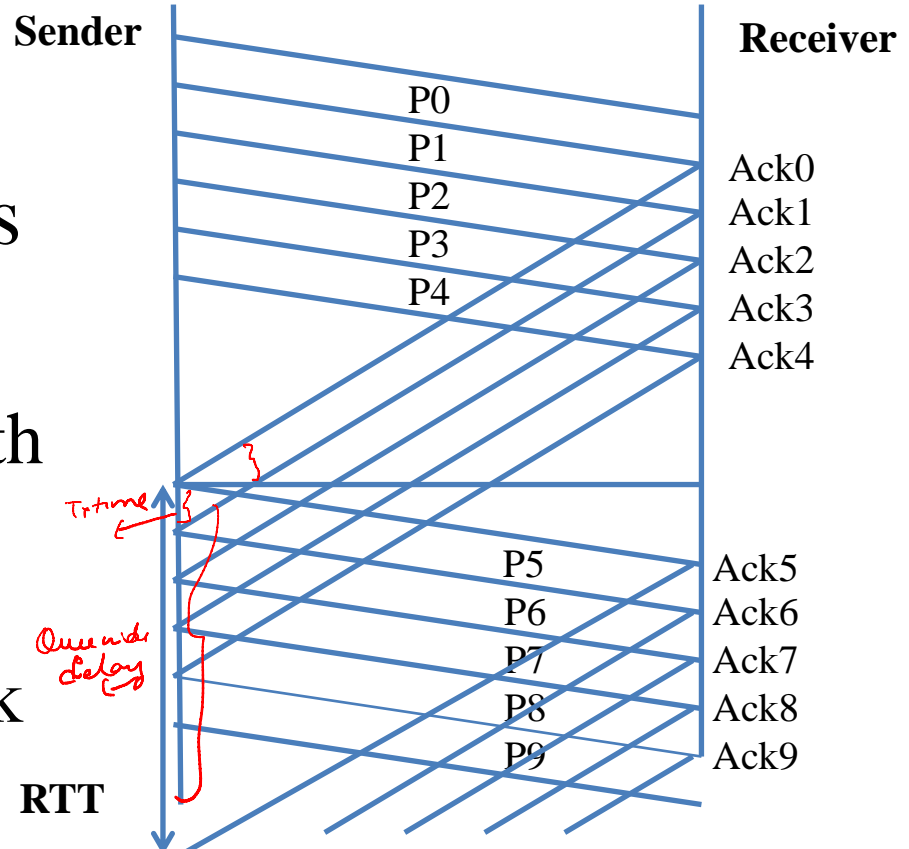
Ack7

Ack8


Ack9



- $W > \text{BDP}$ (in bits)
 $\hookrightarrow \text{RTT}_{\text{no load}}$
- Sliding Window says
 - $\text{Thr} = \frac{W}{\text{RTT}}$ $\hookrightarrow \text{RTT}_{\text{load}}$
Bottleneck bandwidth
- Self-clocking says
 - Max thr is bottleneck bandwidth



Window Size



Same Idea

- View network as a pipe
- Estimate Bandwidth-delay product (capacity) dynamically
 - Uses the variable Congestion Window (CW) to track it
- Self Clocking: Captures bottleneck bandwidth
 - Use ACKs to clock out data
 - Not perfect (With competing traffic ack spacing will not be preserved)

3 Steps

- Getting to Equilibrium[↗]
- Conservation at equilibrium_→
 - Don't put new packet unless old one is removed
- Adapting to Path Dynamics

Summary

- Congestion Control is a complex problem
- Need to implement it in the context of the sliding window protocol
 - Self clocking[✓] is a useful feature (we will rely on this to capture bottleneck bandwidth)
 - Need to determine and adapt W (window size) such that you don't underutilize bandwidth or congest the network
- Ahead: Actual details